Analysis of Massive Data Sets

http://www.fer.hr/predmet/avsp

Prof. dr. sc. Siniša Srbljić

Izv. prof. dr. sc. Goran Delač

Izv. prof. dr. sc. Marin Šilić

Izv. prof. dr. sc. Klemo Vladimir

Faculty of Electrical Engineering and Computing
Consumer Computing Laboratory

Clustering

Izv. prof. dr. sc. Goran Delač

Copyright © 2020 CCL: Analysis of massive data sets

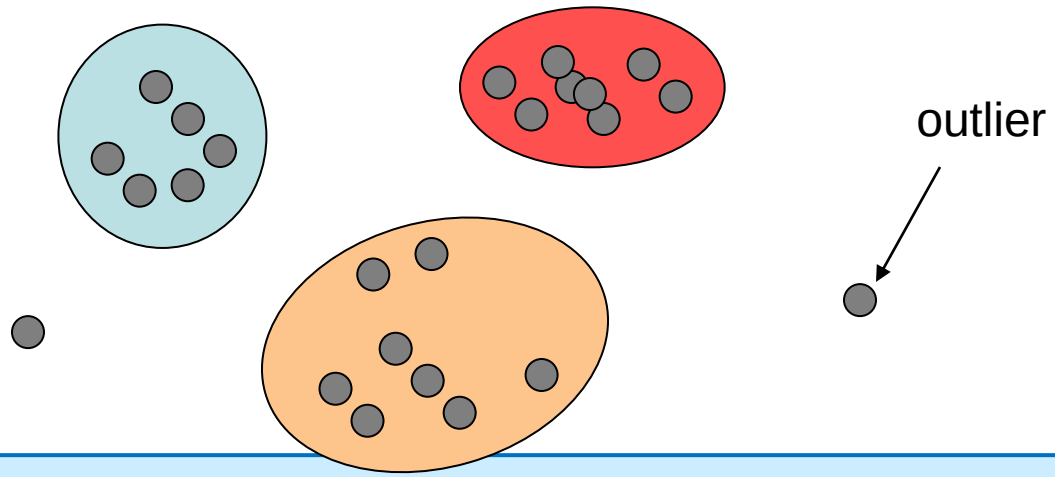
Outline

- **Problem of clustering**
  - Cluster representation
  - Traditional approaches
- **BFR**
- **CURE**

Copyright © 2020 CCL: Analysis of massive data sets

The problem of clustering

- **Why clustering?**
  - To get a better understanding of the data set
  - Given a set of data points, we would like to understand their **structure** – to see which points are **most similar** when compared to other points in the data set



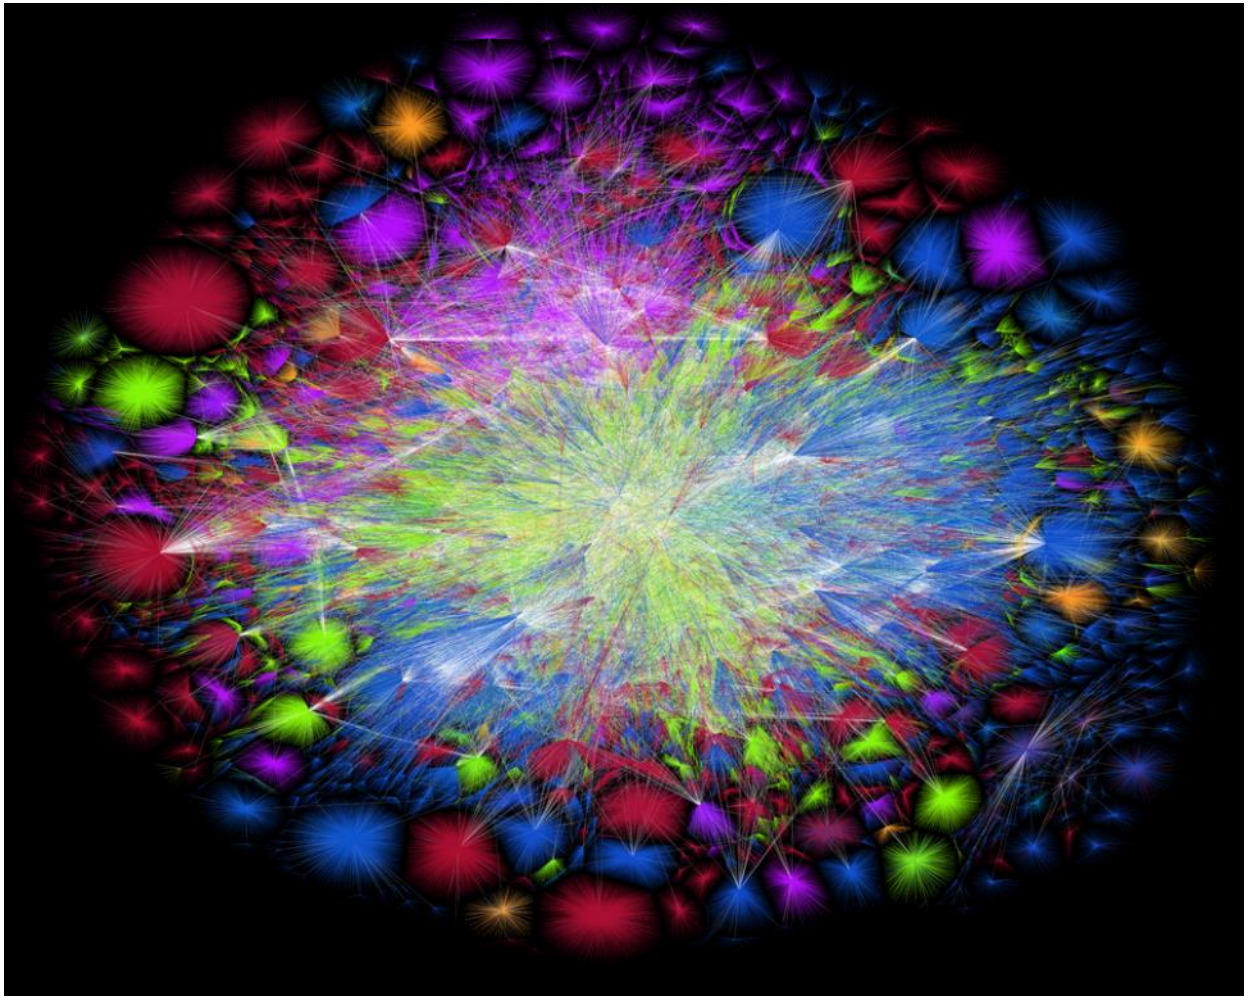
Copyright © 2020 CCL: Analysis of massive data sets

The problem of clustering

- **What is clustering?**
  - A set of points is **grouped** into clusters with a **notion of distance** between the points
    - Goal: group together points that are close to each other (**similar!**)
    - Points in different clusters are distant / not similar
  - Distance
    - Euclidian, Cosine distance, Manhattan distance, Jaccard similarity, Edit distance, Hamming distance, … [2]

Copyright © 2020 CCL: Analysis of massive data sets

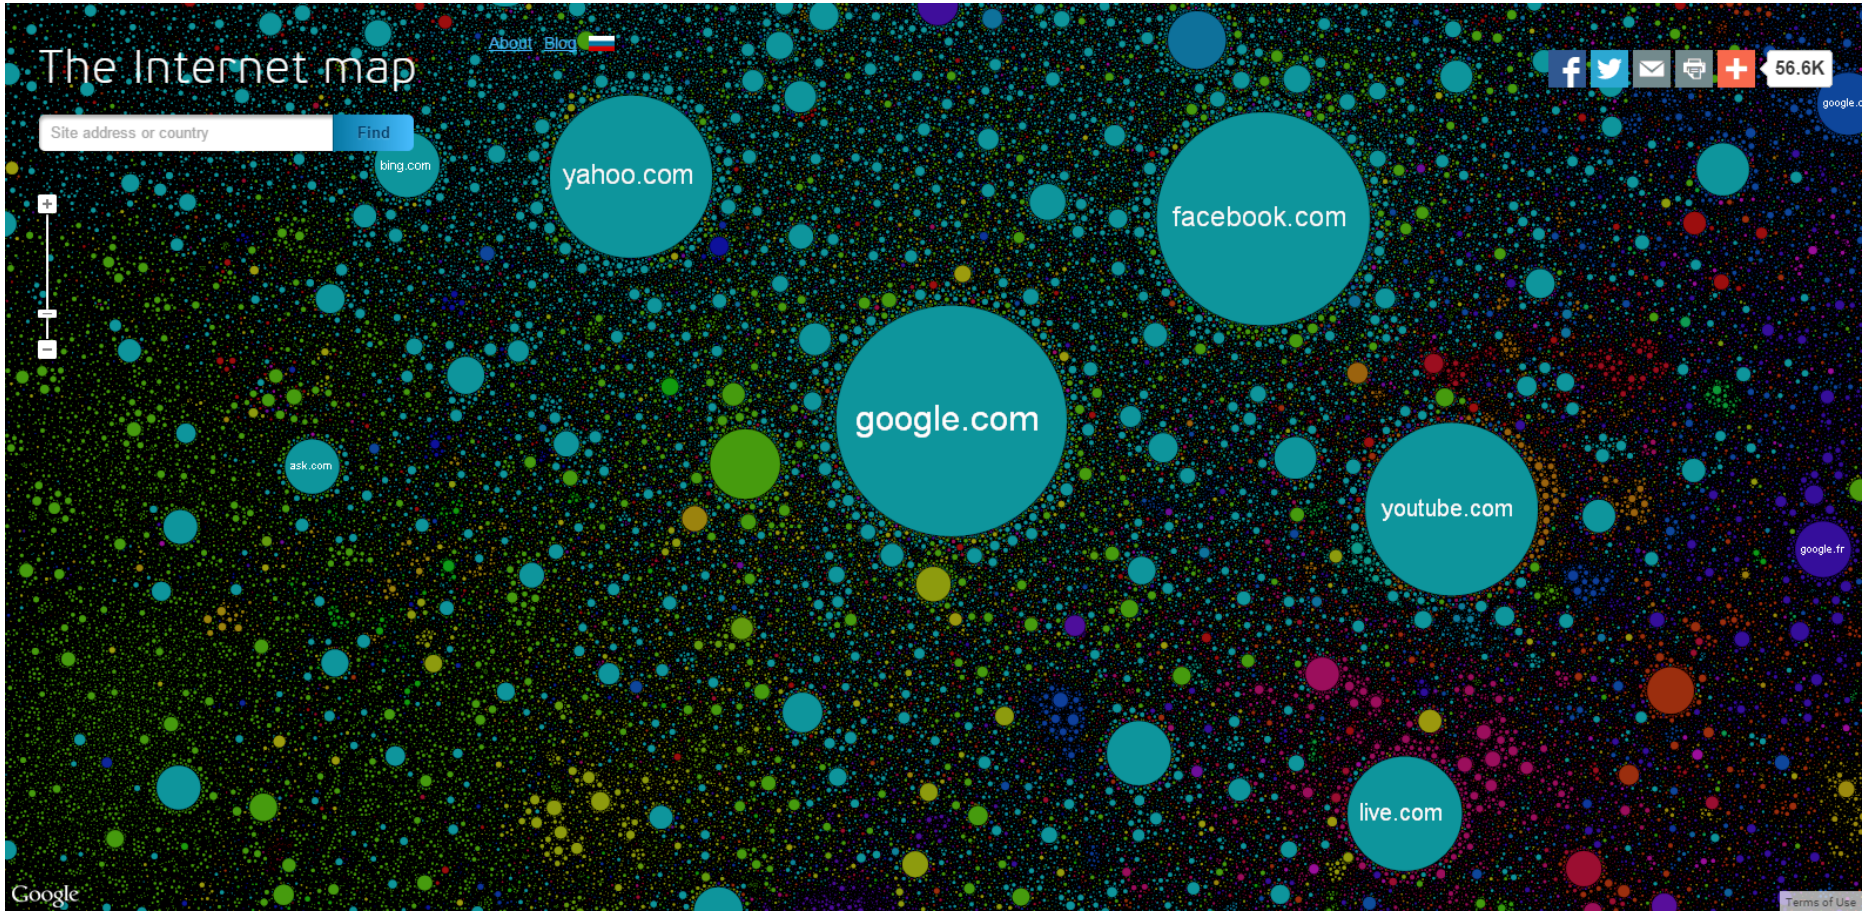
The problem of clustering



The Opte project, https://en.wikipedia.org/wiki/Opte_Project

Copyright © 2020 CCL: Analysis of massive data sets

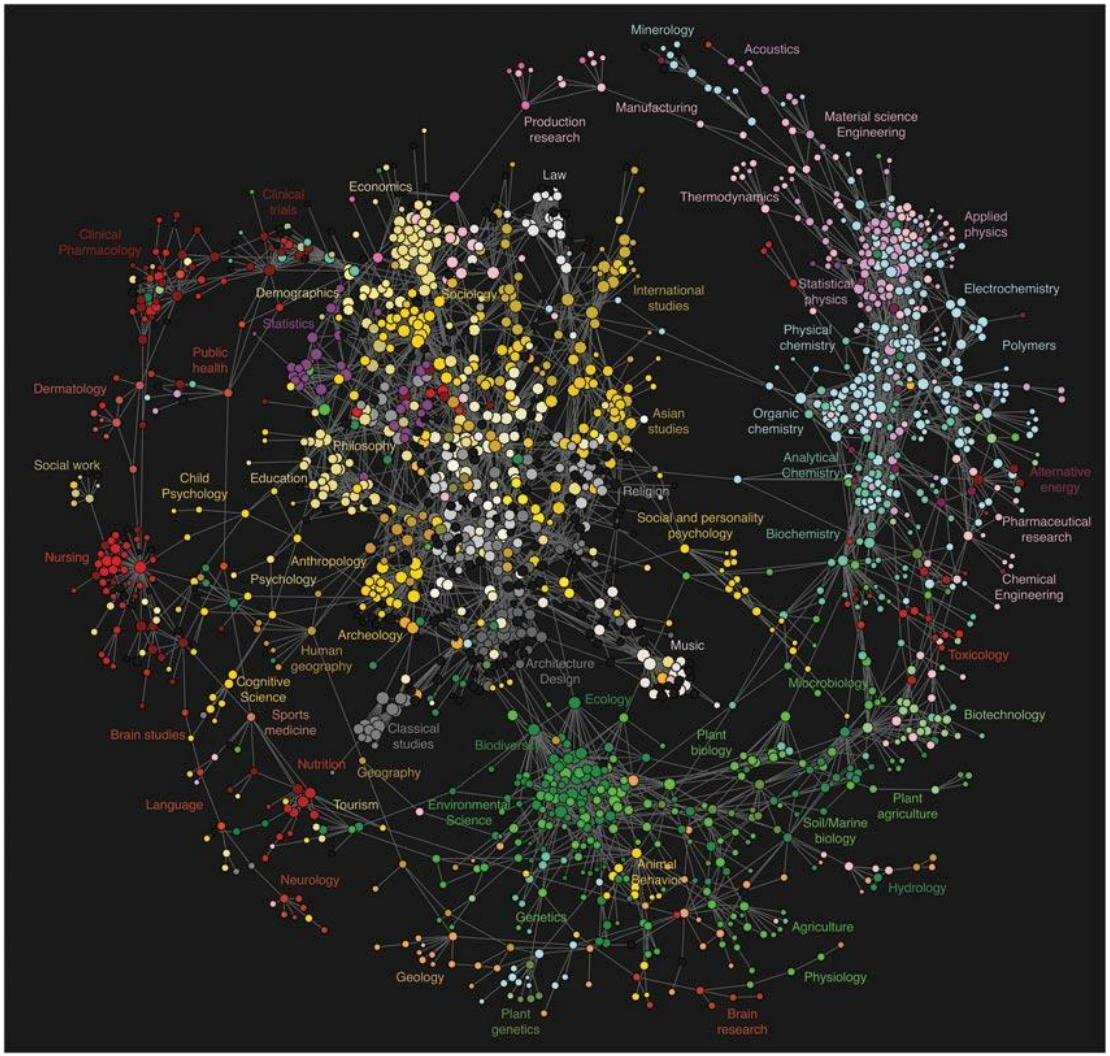
The problem of clustering



The Internet Map, http://internet-map.net/

Copyright © 2020 CCL: Analysis of massive data sets

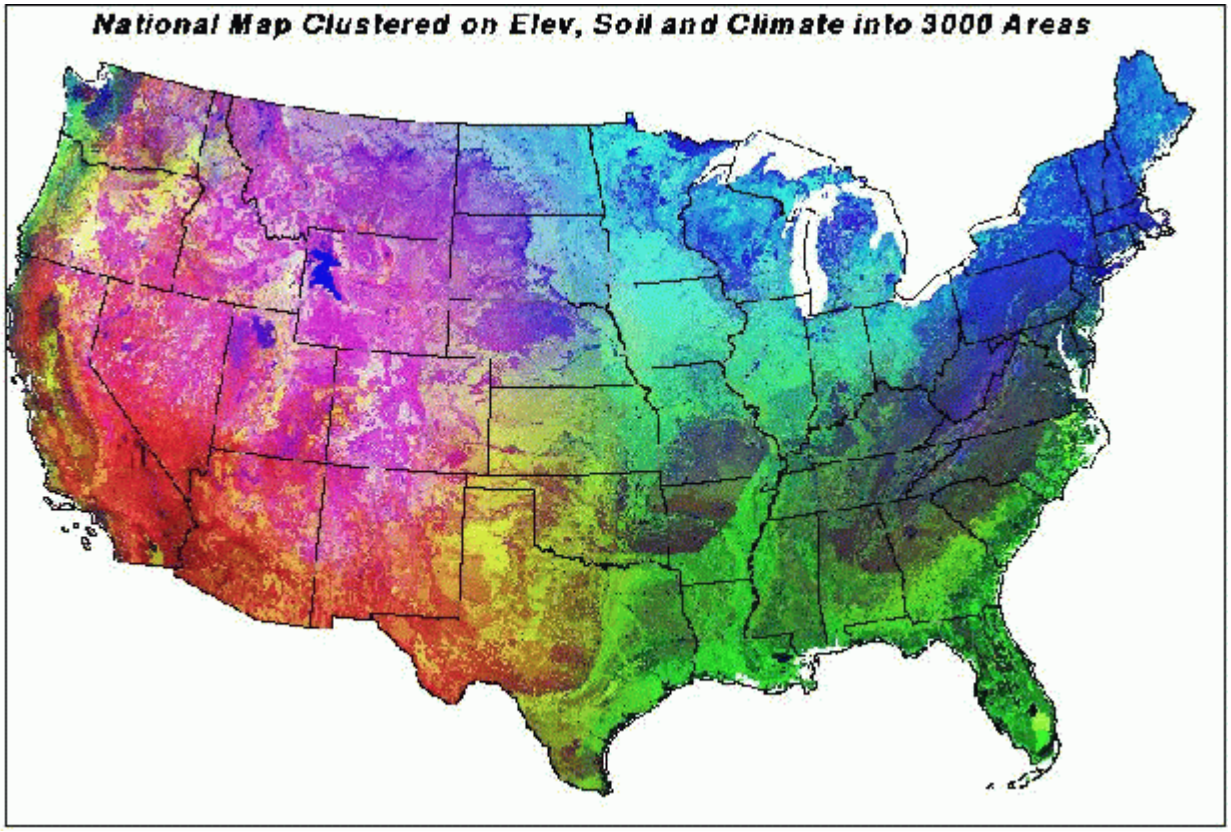
The problem of clustering



Web usage data outline map of knowledge, http://www.nature.com/news/2009/090309/full/458135a.html

Copyright © 2020 CCL: Analysis of massive data sets

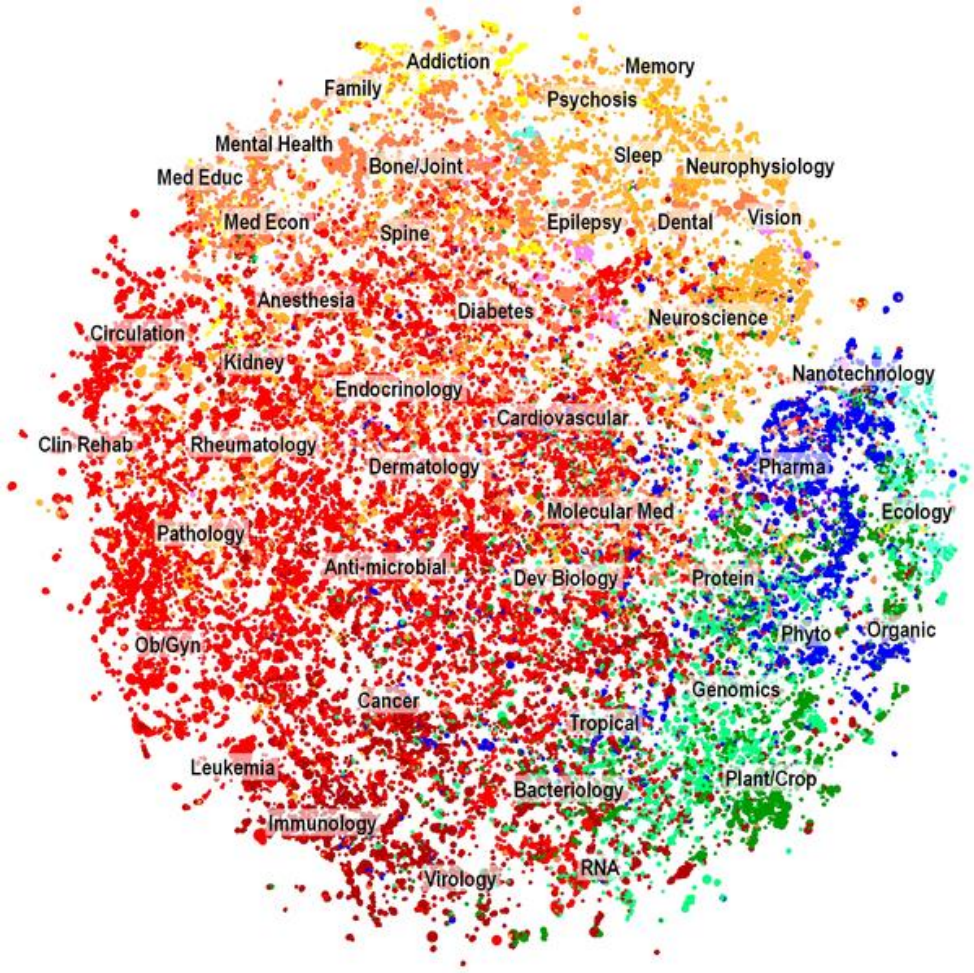
The problem of clustering



National Map of Vegetation Ecoregions Produced Empirically Using Multivariate Spatial Clustering, http://www.geobabble.org/~hnw/esri98/

Copyright © 2020 CCL: Analysis of massive data sets

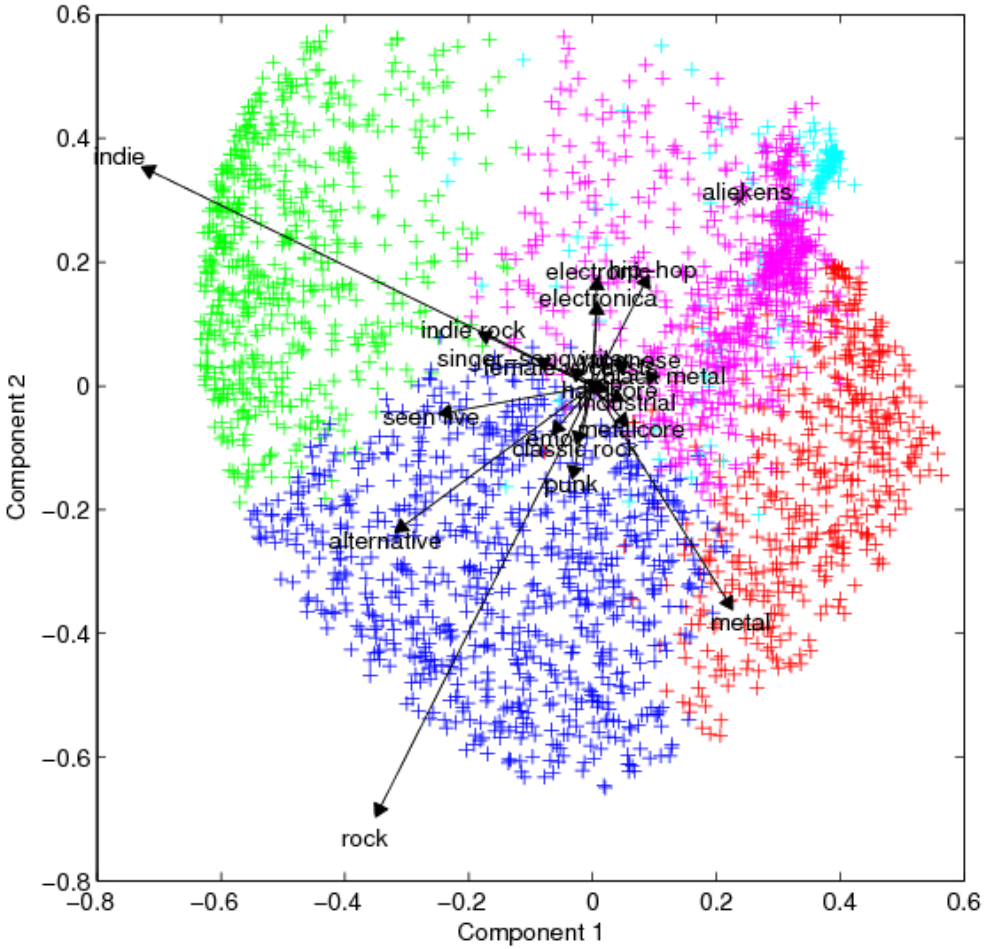
The problem of clustering



Clustering Map of Biomedical Articles (text similarity),
https://journals.plos.org/plosone/article/figure?id=10.1371/journal.pone.0018029.g004

Copyright © 2020 CCL: Analysis of massive data sets

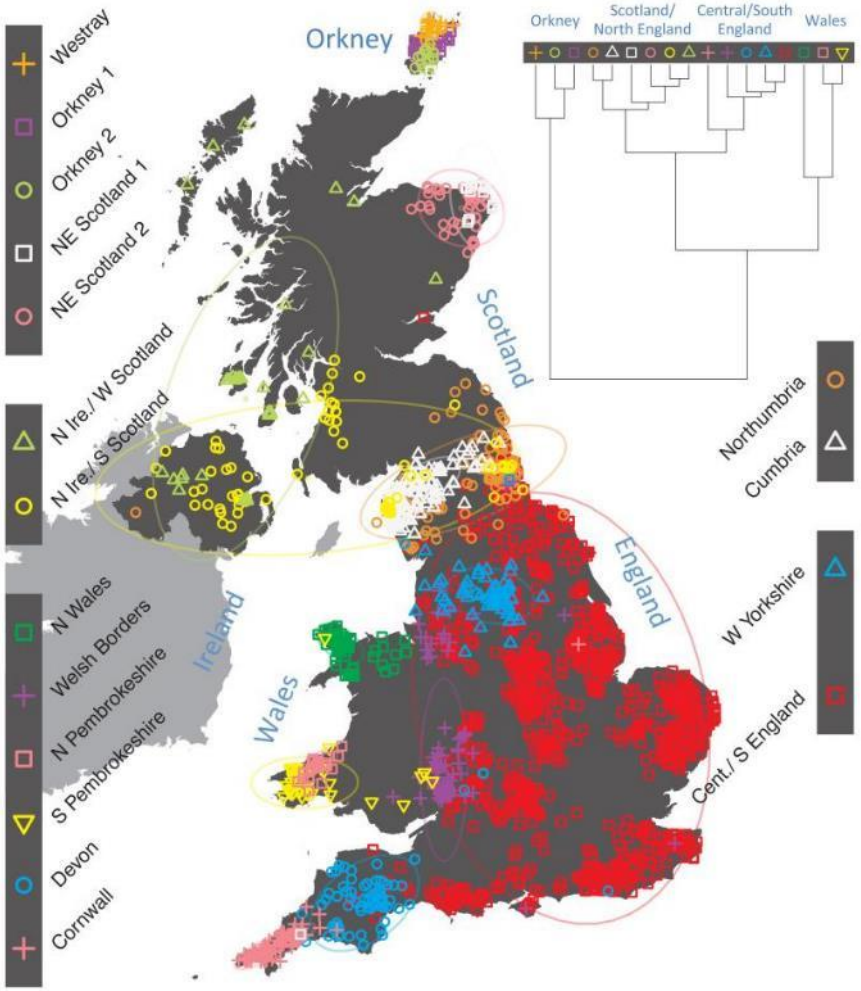
The problem of clustering



Music profiles (Last.fm),
http://anthony.liekens.net/index.php/Computers/DataMining

Copyright © 2020 CCL: Analysis of massive data sets

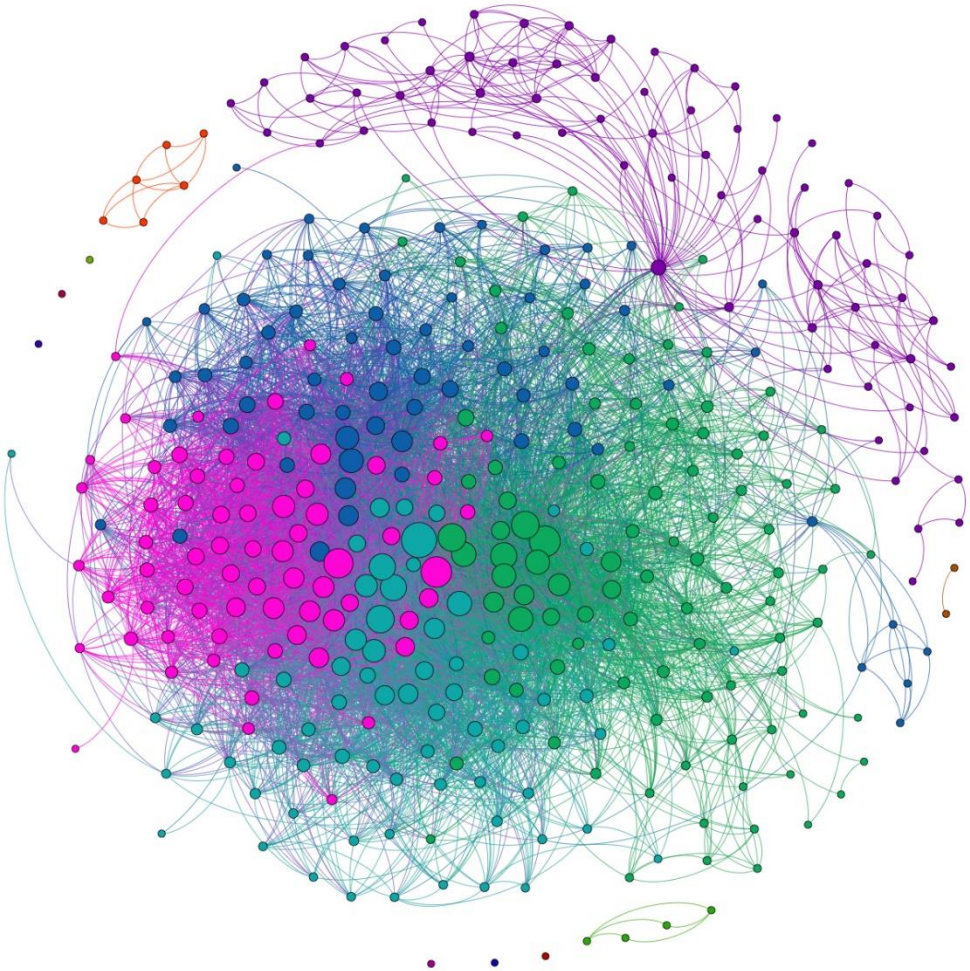
The problem of clustering



Clustering of the 2,039 UK individuals into 17 clusters based only on genetic data, http://www.ncbi.nlm.nih.gov/pmc/articles/PMC4632200/

Copyright © 2020 CCL: Analysis of massive data sets

The problem of clustering



Social network graphs,
http://social-dynamics.org/tag/clustering-algorithm/

Copyright © 2020 CCL: Analysis of massive data sets

The problem of clustering

- **Hard problem!**
  - High-dimensional data
  - A large amount of data points (possibly does not fit into main memory)

- **Curse of high dimensionality**
  - In a high-dimensional space almost all pairs of data points are equally distant from each other

Copyright © 2020 CCL: Analysis of massive data sets

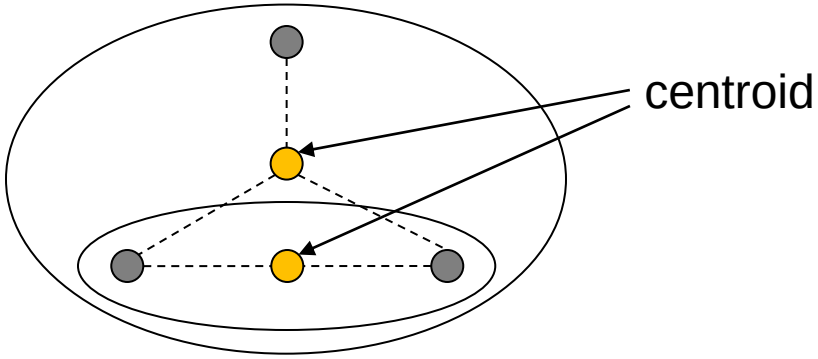
Cluster representation

- **How to represent multiple points?**
  - In **Euclidian** space, clusters can be represented by **centroids**

$$c_i = \frac{\sum_{j=1}^{N} x_{i,j}}{N}$$

$1 \leq i \leq D$ dimension

$1 \leq j \leq N$ point index



Copyright © 2020 CCL: Analysis of massive data sets

Cluster representation

- **How to represent multiple points?**
  - In **Euclidian** space, clusters can be represented by **centroids**

Distance to cluster = distance to centroid

Copyright © 2020 CCL: Analysis of massive data sets

Cluster representation

- **Non-Euclidian spaces**
  - Not possible to compute centroid
  - **1) How to represent a cluster?**
  - **2) How to define "nearness" of clusters**

1. **Clusteroid** – actual representative point is chosen (a point from the data set)
   - Treat clusteroid as a cluster with an appropriate notion of proximity (Jaccard distance, Edit distance, …)

How to pick a clusteroid?

- Point that is ***"closest"*** to all other points in cluster
  - Smallest maximum distance to other points
  - Smallest average distance to other points
  - Smallest sum of squares of distances to other points

Copyright © 2020 CCL: Analysis of massive data sets

Cluster representation

2. Cluster Cohesion – how close are clusters

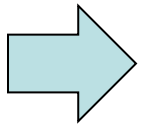
- **Diameter** of a merged cluster
  - Maximum distance between two points
- **Average distance**
  - Between all the points in the cluster
- **Density-based clustering**
  - Number of points in a measure of volume
  - Number of points divided by diameter or avg. distance

Copyright © 2020 CCL: Analysis of massive data sets

Hierarchical clustering

- Agglomerative (bottom up)
  - In the beginning each point is its own cluster
  - Iteratively merge closest clusters

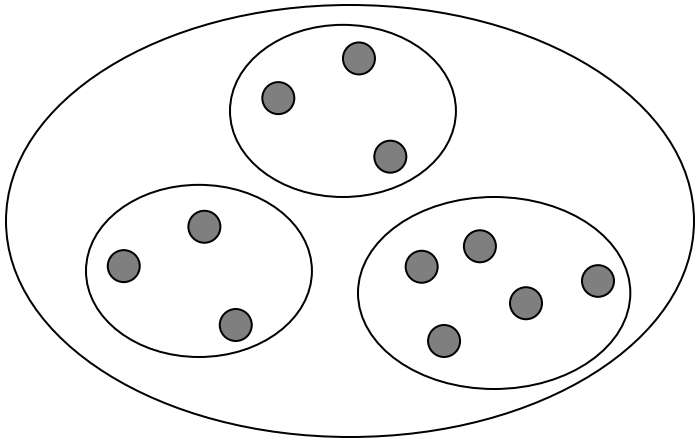
$O(n^3)$



- Divisive (top down)
  - Start with a single cluster
  - Split it recursively

Careful implementation

$O(n^2 \log(n))$



Too slow for large datasets!

Copyright © 2020 CCL: Analysis of massive data sets

Point assignment methods

- **Maintain a set clusters**
  - Known number of clusters
- **Points are assigned to their nearest cluster**
- **K-means clustering**
- **Better suited for large data sets**

Copyright © 2020 CCL: Analysis of massive data sets

K-means clustering

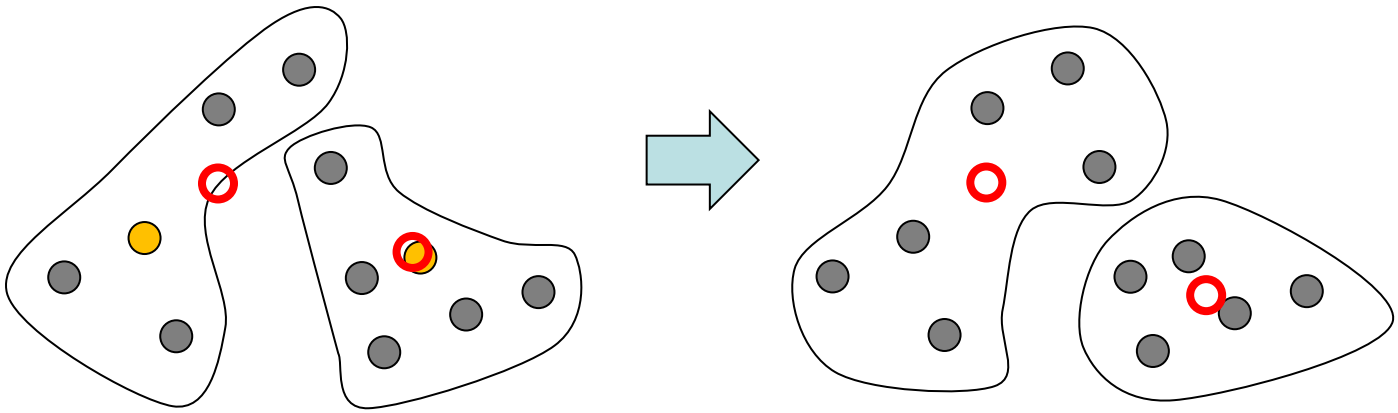
- **Euclidian space**

1. Initialization
   - Choose **k** points to be the initial centroids
   - Assign points to their nearest centroid
2. Recompute centroids' locations
3. Reassign points to new centroids
   - Points can move across cluster boundaries

Copyright © 2020 CCL: Analysis of massive data sets

K-means clustering

- Steps 2-3 are repeated until convergence
  - Points do not change clusters
  - OR some other stopping condition is met



Copyright © 2020 CCL: Analysis of massive data sets

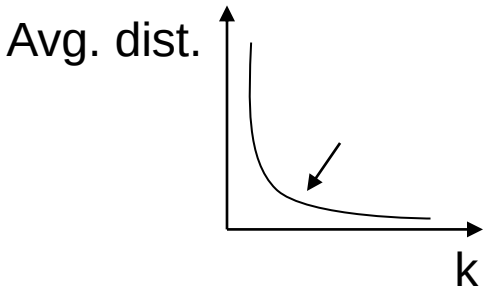
K-means clustering

- How to choose k?
- How to choose k points from the data set?

Copyright © 2020 CCL: Analysis of massive data sets

K-means clustering

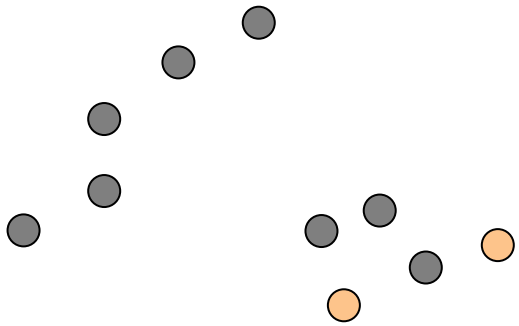
- **Choosing k**
  - Intuitively – if we know the number of clusters in advance or we want to cluster the data in a specific number of clusters
  - Trying different values and observing cluster quality
    - E.g. increasing k will decrease average distance to the centroid within a cluster



Copyright © 2020 CCL: Analysis of massive data sets

K-means clustering

- **Choosing k points**
  - At random? Not a good idea



Copyright © 2020 CCL: Analysis of massive data sets

K-means clustering

- **Choosing k points**
  - Sampling
    - Get a representative sample of the data set
    - Cluster it using hierarchical clustering to get **k** clusters
    - Choose **k** points closest to the centroids
  - Disperse points
    - Pick first point at random
    - Choose the next point to be the most distant one to the selected points (the one with the greatest minimal distance)
    - Repeat the process until k points are selected

Copyright © 2020 CCL: Analysis of massive data sets

K-means clustering

- **Computational complexity**

$$O(rkN)$$

k – number of centroids
N – number of points
r – number of rounds

  - Number of rounds (r) can get large
  - N can be really large in massive data sets

  - Can we get meaningful clustering results in just one pass?

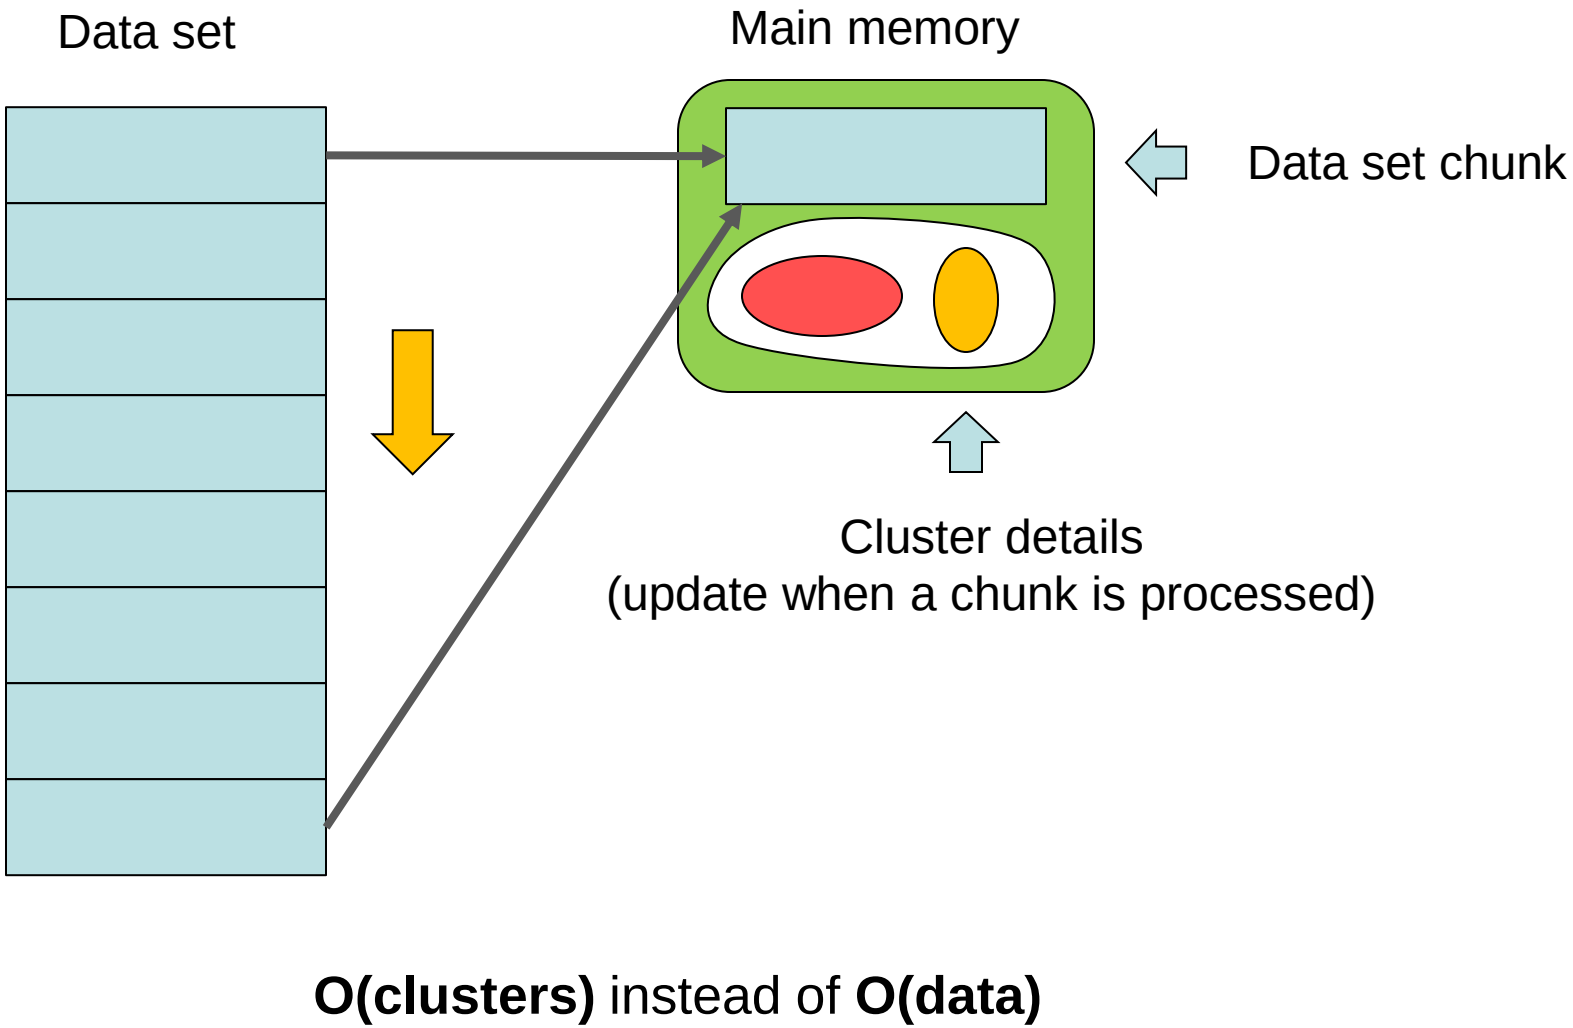
Copyright © 2020 CCL: Analysis of massive data sets

BFR

- **Bradley, Fayyad, Reina**
- **Based on k-means clustering approach**
- **Idea**
  - Avoid storing information on all points as they do not fit into main memory
  - Keep track of summarized cluster information

Copyright © 2020 CCL: Analysis of massive data sets

BFR

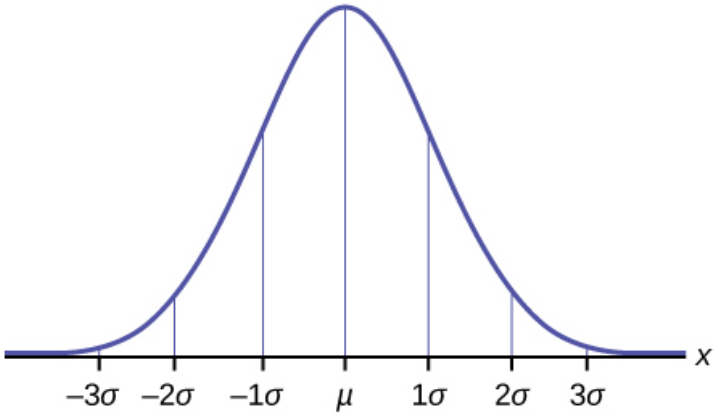


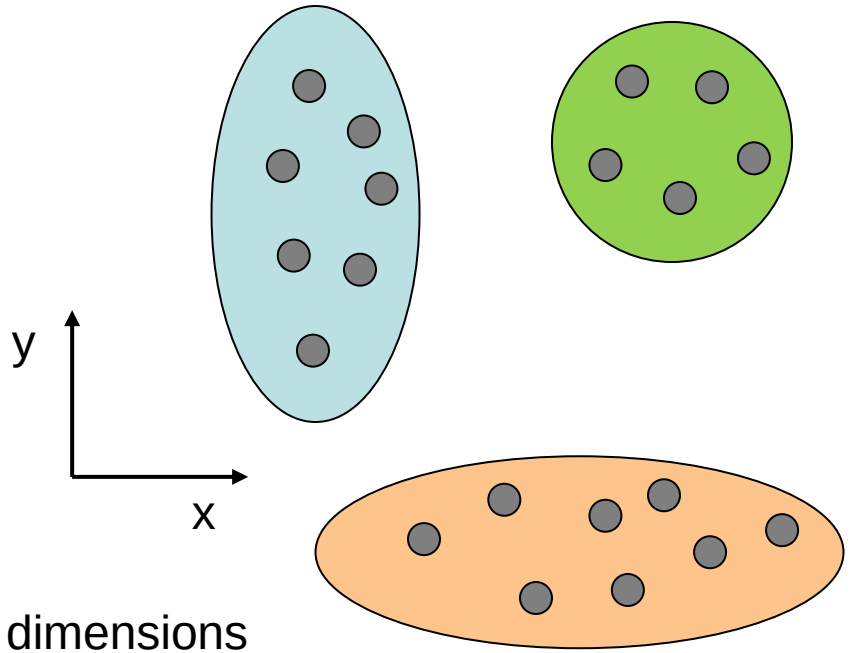
O(clusters) instead of **O(data)**

Copyright © 2020 CCL: Analysis of massive data sets

BFR

- **Assumptions**
  - Cluster members are normally distributed around the centroid
  - Euclidian space



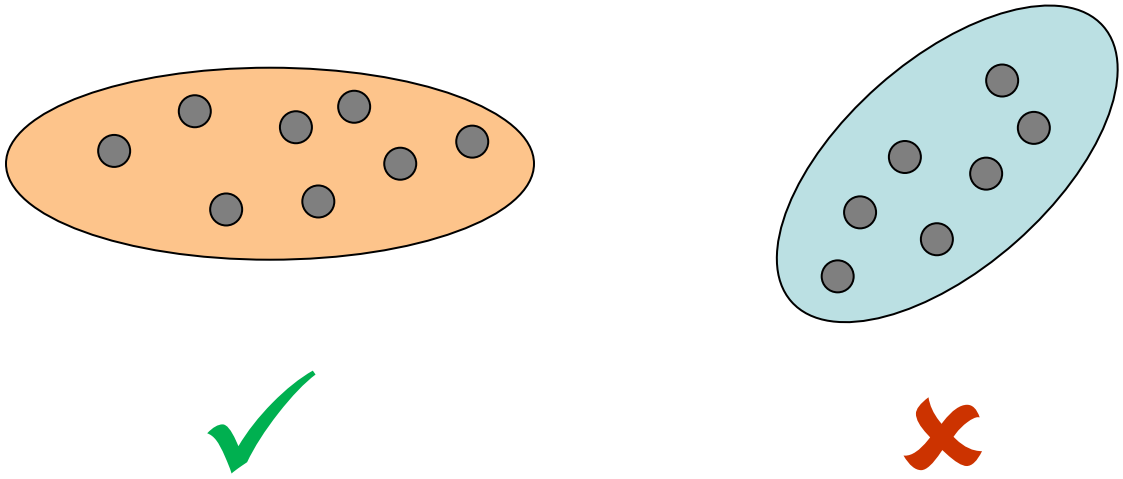


Standard deviations can vary for different dimensions

Copyright © 2020 CCL: Analysis of massive data sets

BFR

- **Assumptions**
  - Clusters are axis-aligned ellipses



Copyright © 2020 CCL: Analysis of massive data sets

BFR

Algorithm

1. Choose initial centroids
   - Some appropriate method, like in k-means clustering
2. Process data set
   - Read data set chunk at a time
   - Maintain cluster information
3. Finalize cluster information

Copyright © 2020 CCL: Analysis of massive data sets

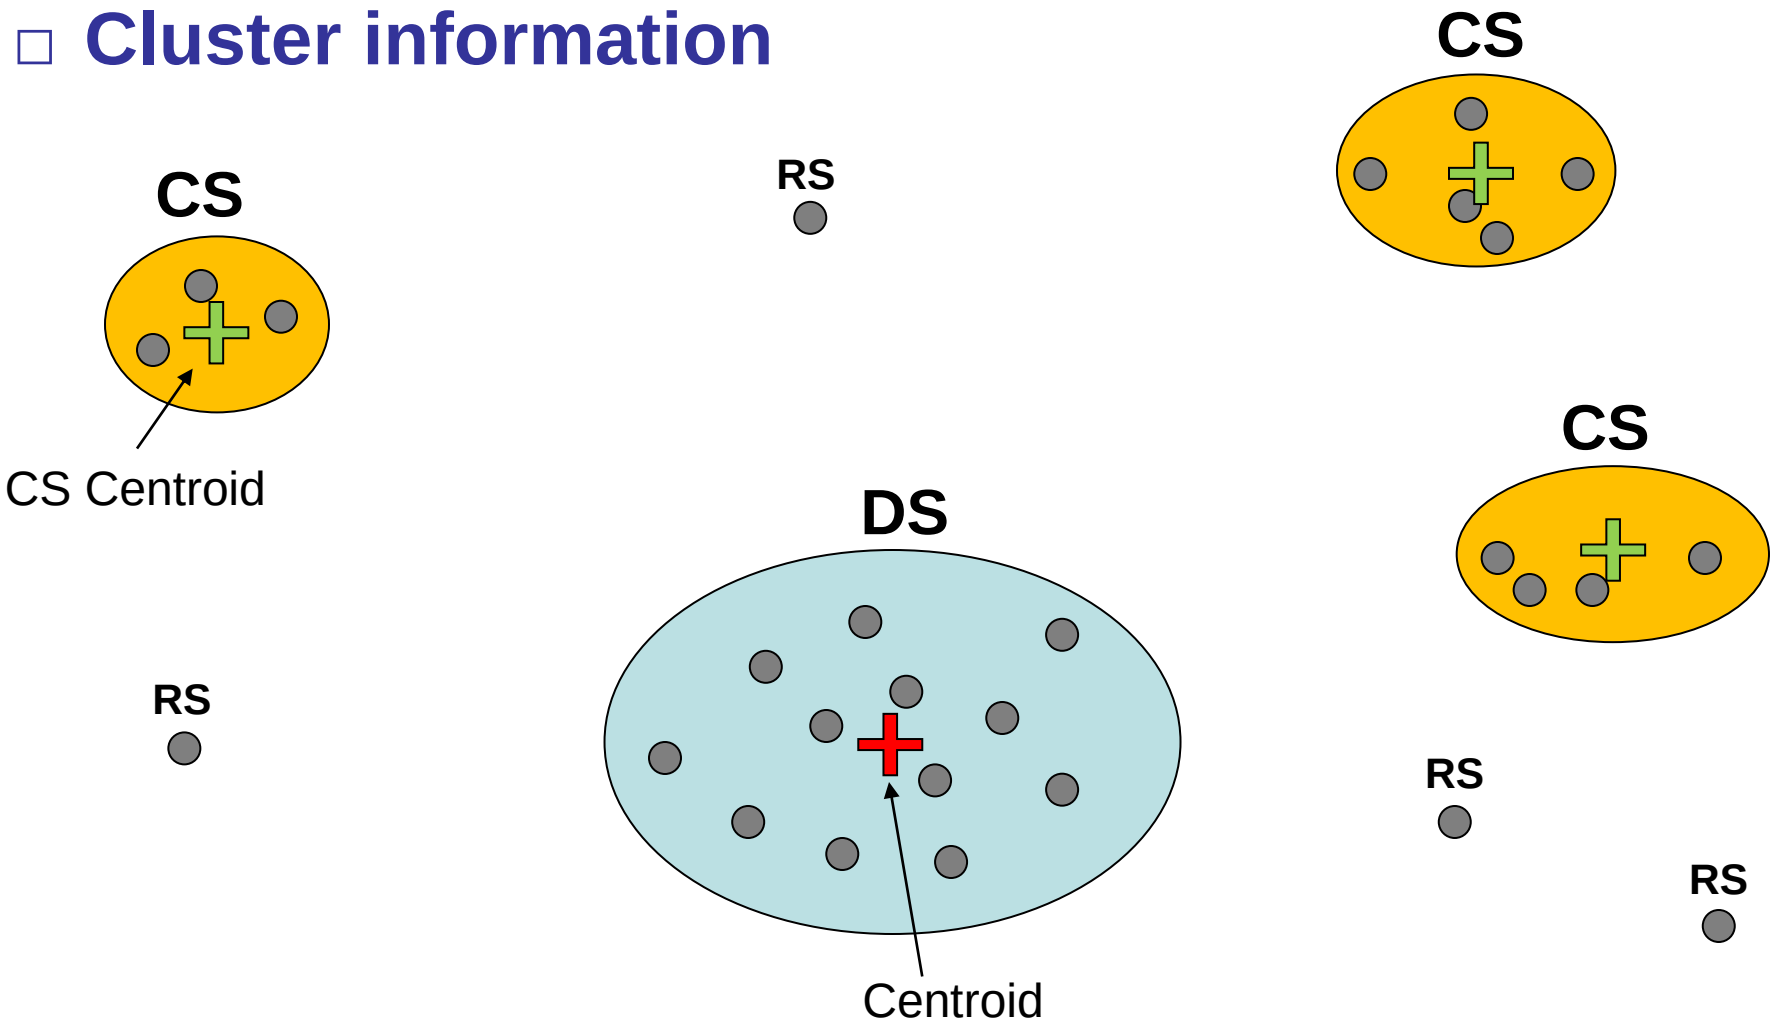
BFR

- **Cluster information**
  - Discard set (**DS**)
    - Points that are close enough to the centroid to be summarized and then discarded
  - Compression set (**CS**)
    - Groups of points that are close together but not close to any cluster centroid
    - Points are summarized in a compression set and discarded (the are not assigned to any cluster!)
  - Retained set (**RS**)
    - Isolated points – kept in memory

Copyright © 2020 CCL: Analysis of massive data sets

BFR

- **Cluster information**



Copyright © 2020 CCL: Analysis of massive data sets

BFR

- **Cluster information (DS + CS)**
  - The number of points (**N**)
    - Points that exist in a cluster
  - Vector SUM
    - Each element is a sum of point coordinates across a dimension

$$v_i = \sum_{j=1}^{N} p_{i,j}$$

  - Vector SUMSQ
    - Each element is a sum of squares of point coordinates across a dimension

$$v_i = \sum_{j=1}^{N} p_{i,j}^2$$

Copyright © 2020 CCL: Analysis of massive data sets

BFR

- **Cluster information (DS + CS)**
  - 2d + 1 values are stored to represent each cluster / compression set
    - d is the number of dimensions
  - Centroid
    - Can be calculated as $\mathbf{c_i = SUM_i / N}$ **(not stored!)**
    - $\mathbf{SUM_i}$ sum of coordinates across $\mathbf{i^{th}}$ dimension
  - Variance
    - $\mathbf{v_i = (SUMSQ_i / N) - (SUM / N)^2}$ **(not stored!)**
    - Standard deviation is the square root of $\mathbf{v_i}$

Copyright © 2020 CCL: Analysis of massive data sets

BFR

Cluster information (DS + CS) – Benefits

- (*ds* + *cs*) * (2d + 1) + *rs***d* values are stored at any time
  - *ds* number of clusters **- fixed**
  - *cs* number of compression sets
  - *rs* size of retains set
  - As *cs* increases, *rs* decreases
- Large dataset
  - (ds + cs) * (2d + 1) + rs * d << N * d

Copyright © 2020 CCL: Analysis of massive data sets

BFR

□ Cluster information (DS + CS) – Benefits

- It is easier to compute new centroid from SUM_i
  - Updates are simple additions to existing vector
- Much easier to compute new variance using SUM_i and $SUMSQ_i$
- Very efficient when combining sets
  - Simple computation of a new centroid and variance

Copyright © 2020 CCL: Analysis of massive data sets

BFR

□ Dataset chunk processing

1. Load chunk from memory

2. Find points that are ***close enough*** to a cluster centroid
   - Update cluster data
   - Discard points

BFR

□ Dataset chunk processing

3. Cluster the remaining points along with the points in the retain set **(RS)**
   - Points are added to compression sets **(CS)**
   - Any main-memory cluster algorithm can be applied
   - Define criterion when to add points to clusters and when to merge clusters
   - CS data is updated and points are discarded!

4. Check if some compression sets **(CS)** can be merged

Copyright © 2020 CCL: Analysis of massive data sets

BFR

- **Finalized cluster information**
  - CS and RS sets can be treated as outliers and discarded
  - CS and RS are merged with a cluster whose centroid is **close enough**

Copyright © 2020 CCL: Analysis of massive data sets

BFR

- When is a point close enough to be included into a cluster?

- When to merge compression sets?

Copyright © 2020 CCL: Analysis of massive data sets

BFR

- **Mahalanobis distance (MD)**
  - If **MD** is **below** a certain **threshold** - likelihood that a point **belongs** to a certain cluster is **high**
  - Normalized Euclidian distance form centroid

$$MD(x,c) = \sqrt{\sum_{i=1}^{d} \left(\frac{x_i - c_i}{\sigma_i}\right)^2}$$

    - x point
    - c centroid of the cluster
    - d number of dimensions
    - σ_i standard deviations of points in the cluster

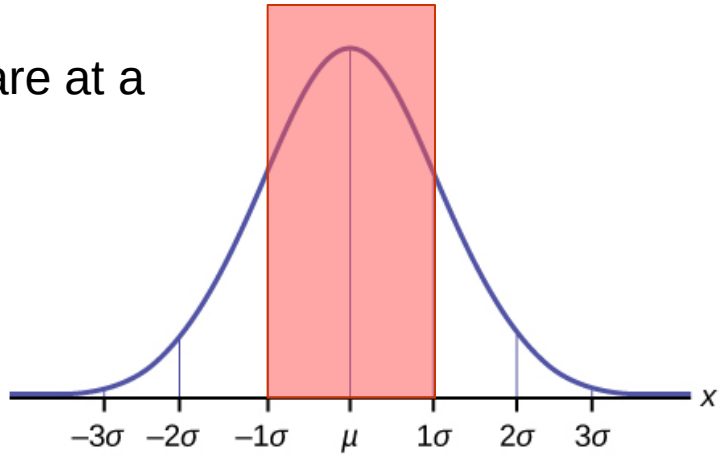
Copyright © 2020 CCL: Analysis of massive data sets

Mahalanobis distance (MD)

- Points are **normally distributed** around the centroid!
- If $x_i - c_i = \sigma_i$

$$MD(x,c) = \sqrt{\sum_{i=1}^{d} \left(\frac{x_i - c_i}{\sigma_i}\right)^2} = \sqrt{\sum_{i=1}^{d} \left(\frac{\sigma_i}{\sigma_i}\right)^2} = \sqrt{d}$$

68% points in the cluster are at a distance of $\sqrt{d}$

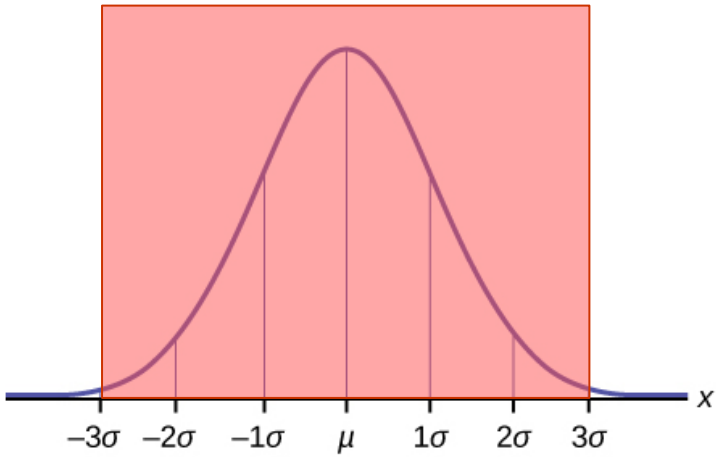


Copyright © 2020 CCL: Analysis of massive data sets

BFR

□ Mahalanobis distance (MD)

> 99% points in the cluster are at a distance of $3\sqrt{d}$



Appropriate threshold can be chosen, e.g. $3\sqrt{d}$

Copyright © 2020 CCL: Analysis of massive data sets

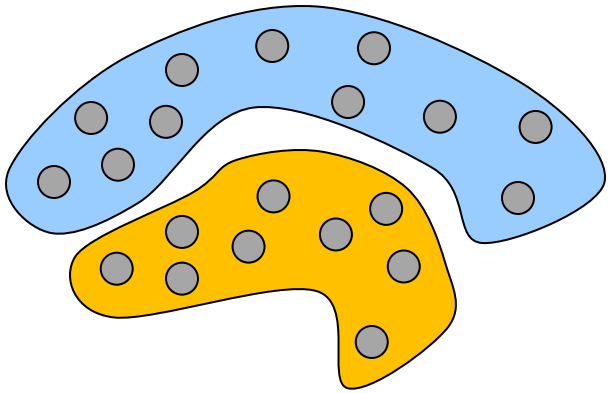
BFR

- **When to merge compression sets?**
  - Calculate combined variance of merged two CS
    - SUM, SUMSQ make this operation easy
  - If variance is below a certain threshold, merge the compression sets
  - Many alternative approaches
    - Treat dimensions differently, as they might differ in importance
    - Cluster density

Copyright © 2020 CCL: Analysis of massive data sets

CURE

- **Clustering using representatives**
  - Can handle massive datasets
  - Euclidian space
  - No assumptions on the cluster shape!



Copyright © 2020 CCL: Analysis of massive data sets

CURE

- **1. pass**

1. Take a data set sample that fits into main memory

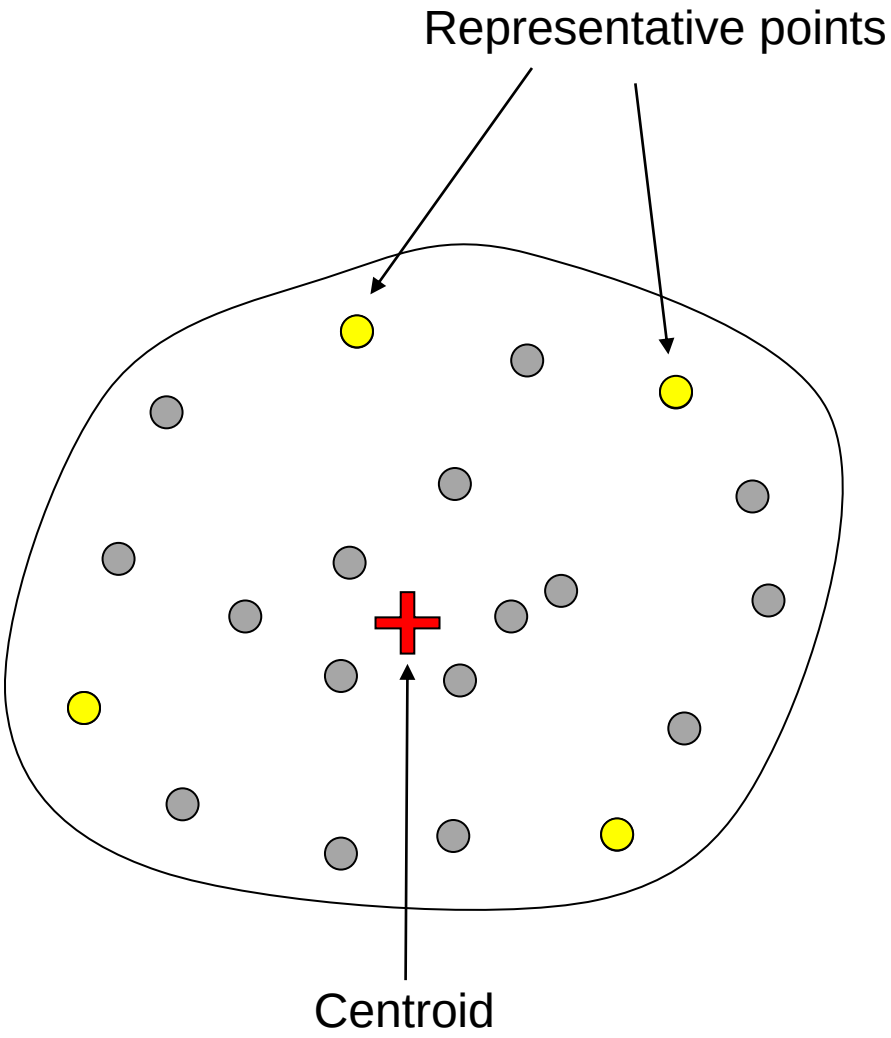
2. Cluster the sample data set
   - Any algorithm that can handle oddly shaped clusters
   - Usually hierarchical clustering

3. Choose **n** representative points
   - **Points should be as distant from each other as possible (see previously described methods)**
   - Usually a small set of points is chosen

Copyright © 2020 CCL: Analysis of massive data sets

CURE

- **1. pass**



Copyright © 2020 CCL: Analysis of massive data sets

CURE

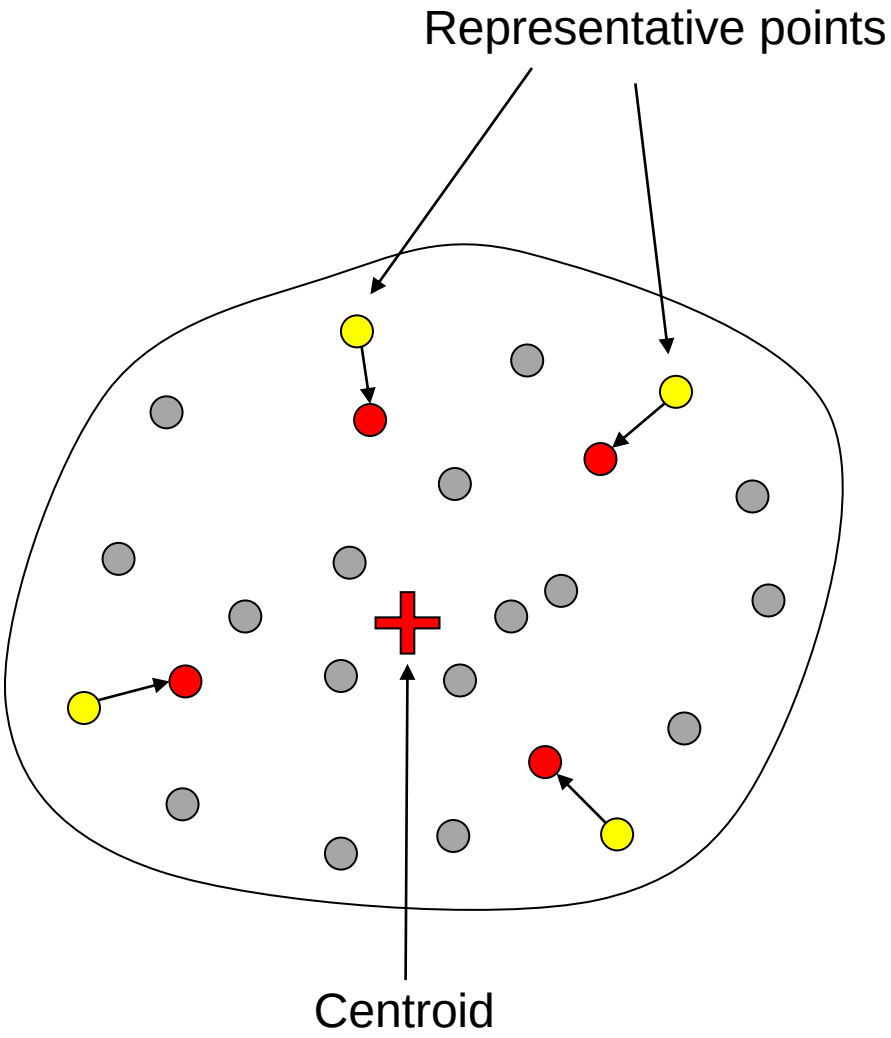
- **1. pass**

4. Move the representative points some portion of a distance towards the centroid
   - e.g. 20%
   - The data set is not actually changed, these are simply synthetic points used by the algorithm!

Copyright © 2020 CCL: Analysis of massive data sets

CURE

- **1. pass**

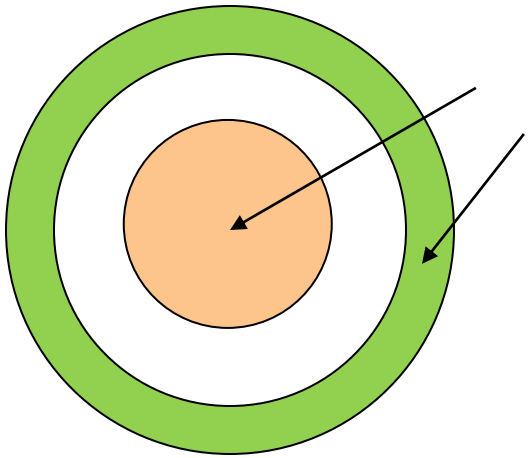


Copyright © 2020 CCL: Analysis of massive data sets

CURE

- **1. pass**

5. Check if clusters can be merged
   - Merge clusters if they have **representative points** that are sufficiently close
   - Note that distance to centroids is not taken into account!



Both clusters have the same centroid!

Copyright © 2020 CCL: Analysis of massive data sets

CURE

- **2. pass**
  - Go through all points **p** in the original data set
  - Assign **p** to a cluster **c** if the point is closest to one of its representative points
    - All other representative points across all clusters are more distant

Copyright © 2020 CCL: Analysis of massive data sets

Literature

1. J. Leskovec, A. Rajaraman, and J. D. Ullman, "Mining of Massive Datasets", 2014, Chapter 7 : “Clustering" (link)
2. Distance and Similarity measures: https://reference.wolfram.com/language/guide/DistanceAndSimilarityMeasures.html

Copyright © 2020 CCL: Analysis of massive data sets